



✧ The CNN challenge

Deep Learning

✧
Jelena Duric
Esteban Aspe
Nicoleta Roman

Process

01

Tried 5
architectures

02

Tried data
augmentation

03

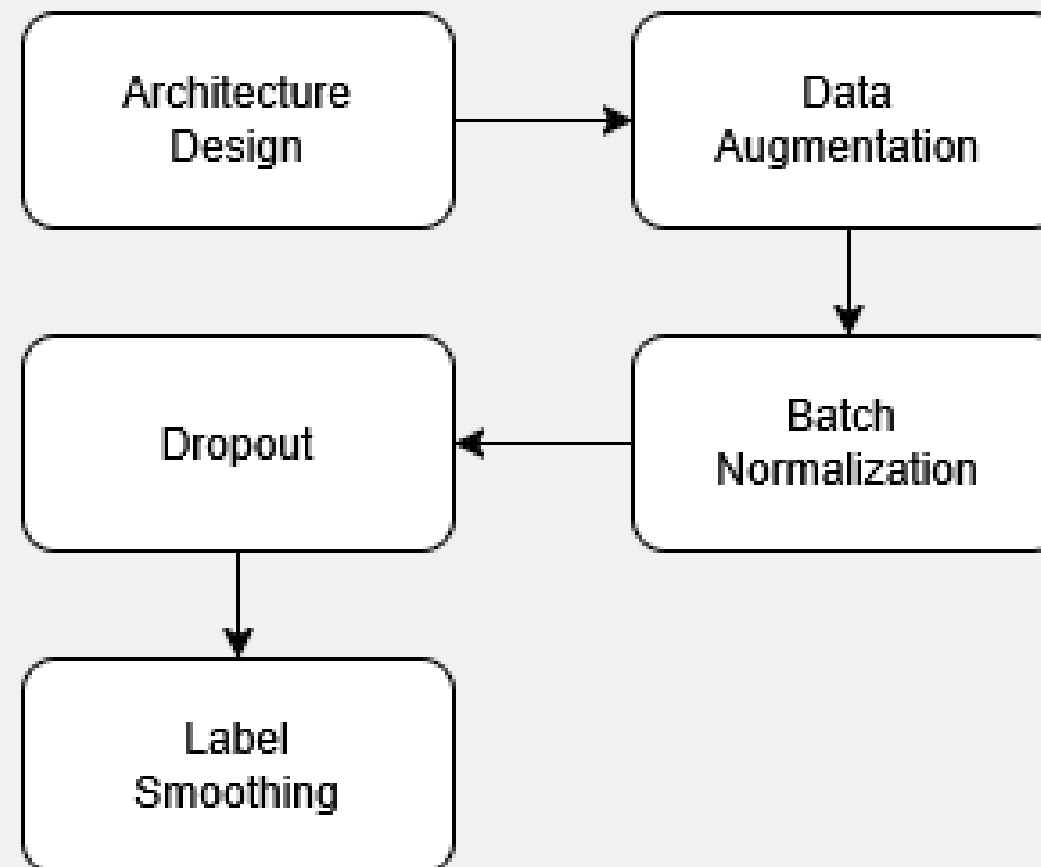
Tried with Batch
Normalization

04

Tried with
different Dropout
rates

05

Tried with label
smoothing





Architecture 1

- Simple CNN

Conv(32 filters, 3×3) + leaky ReLU

Conv(32 filters, 3×3) + leaky ReLU

MaxPooling(2×2)

Conv(64 filters, 3×3) + leaky ReLU

Conv(64 filters, 3×3) + Leaky ReLU

MaxPooling(2×2)

Fully connected (128 neurons) + leaky ReLU



Architecture 2

- **Deeper CNN (VGG-style)**

Conv(64 filters, 3×3) + leaky ReLU

Conv(64 filters, 3×3) + leaky ReLU

MaxPooling(2×2)

Conv(128 filters, 3×3) + leaky ReLU

Conv(128 filters, 3×3) + ReLU

MaxPooling(2×2)

Conv(256 filters, 3×3) + leaky ReLU

Conv(256 filters, 3×3) + leaky ReLU

MaxPooling(2×2)

Fully connected (256 neurons) + leaky ReLU



Architecture 3

- **Residual CNN
(ResNet Inspired,
but smaller)**

Conv(64 filters, 3×3) + leaky ReLU

Residual Block:

Conv(64, 3×3) + leaky ReLU

Conv(64, 3×3) (skip connection)

MaxPooling(2×2)

Conv(128 filters, 3×3) + leaky ReLU

Residual Block:

Conv(128, 3×3) + leaky ReLU

Conv(128, 3×3) (skip connection)

MaxPooling(2×2)

Fully connected (256 neurons) + leaky ReLU



Architecture 4

- Dilated CNN

Conv(64 filters, 3×3 , dilation=2) + leaky ReLU

Conv(64 filters, 3×3 , dilation=2) + leaky ReLU

MaxPooling(2×2)

Conv(128 filters, 3×3 , dilation=2) + leaky ReLU

Conv(128 filters, 3×3 , dilation=2) + leaky ReLU

MaxPooling(2×2)

Fully connected (128 neurons) + leaky ReLU



Architecture 5

- **Extra-Deep CNN (VGG-like)**

Conv layer (64 filters, 3×3) + LeakyReLU

Conv layer (64 filters, 3×3) + LeakyReLU

MaxPooling (2×2)

Conv layer (128 filters, 3×3) + LeakyReLU

Conv layer (128 filters, 3×3) + LeakyReLU

MaxPooling (2×2)

Conv layer (256 filters, 3×3) + LeakyReLU

Conv layer (256 filters, 3×3) + LeakyReLU

MaxPooling (2×2)

Conv layer (512 filters, 3×3) + LeakyReLU

Conv layer (512 filters, 3×3) + LeakyReLU

MaxPooling (2×2)

Dense (256 neurons) + LeakyReLU


Dense (256 neurons) + LeakyReLU

Baseline Results



Architectures used, all of them with leakyReLu and Adam:

- 1.Simple CNN
- 2.Deeper CNN (VGG-style)
- 3.Residual CNN (ResNet-style but smaller)
- 4.Dilated CNN
- 5.Extra-Deep CNN (VGG-style)



Discussion. Data augmentation

- Applied it on training data to architectures 2 and 5
- Applied several image transformations to make the model more flexible and prevent overfitting
- It randomly rotates images up to 20 degrees, shifts them left, right, up, or down by 20%, and slightly tilts (shears) them. It also zooms in and out by 20% and flips images horizontally. If any empty spaces appear due to these changes, they are filled using nearby pixel values
- Kept training on 20 epochs, considering that the overfitting was happening pretty early in the previous network (around third and fourth epoch)
- Results on both network were great, overfitting was almost non existent

Discussion. Batch normalization

- Tried batch normalization on architecture 5 with data augmentation, to prevent overfitting even more
- Used 100 epochs and lowered batch size to 64, considering that with 128 we ran out of memory on the GPU
- Very good results and with no overfitting

Discussion. Dropout

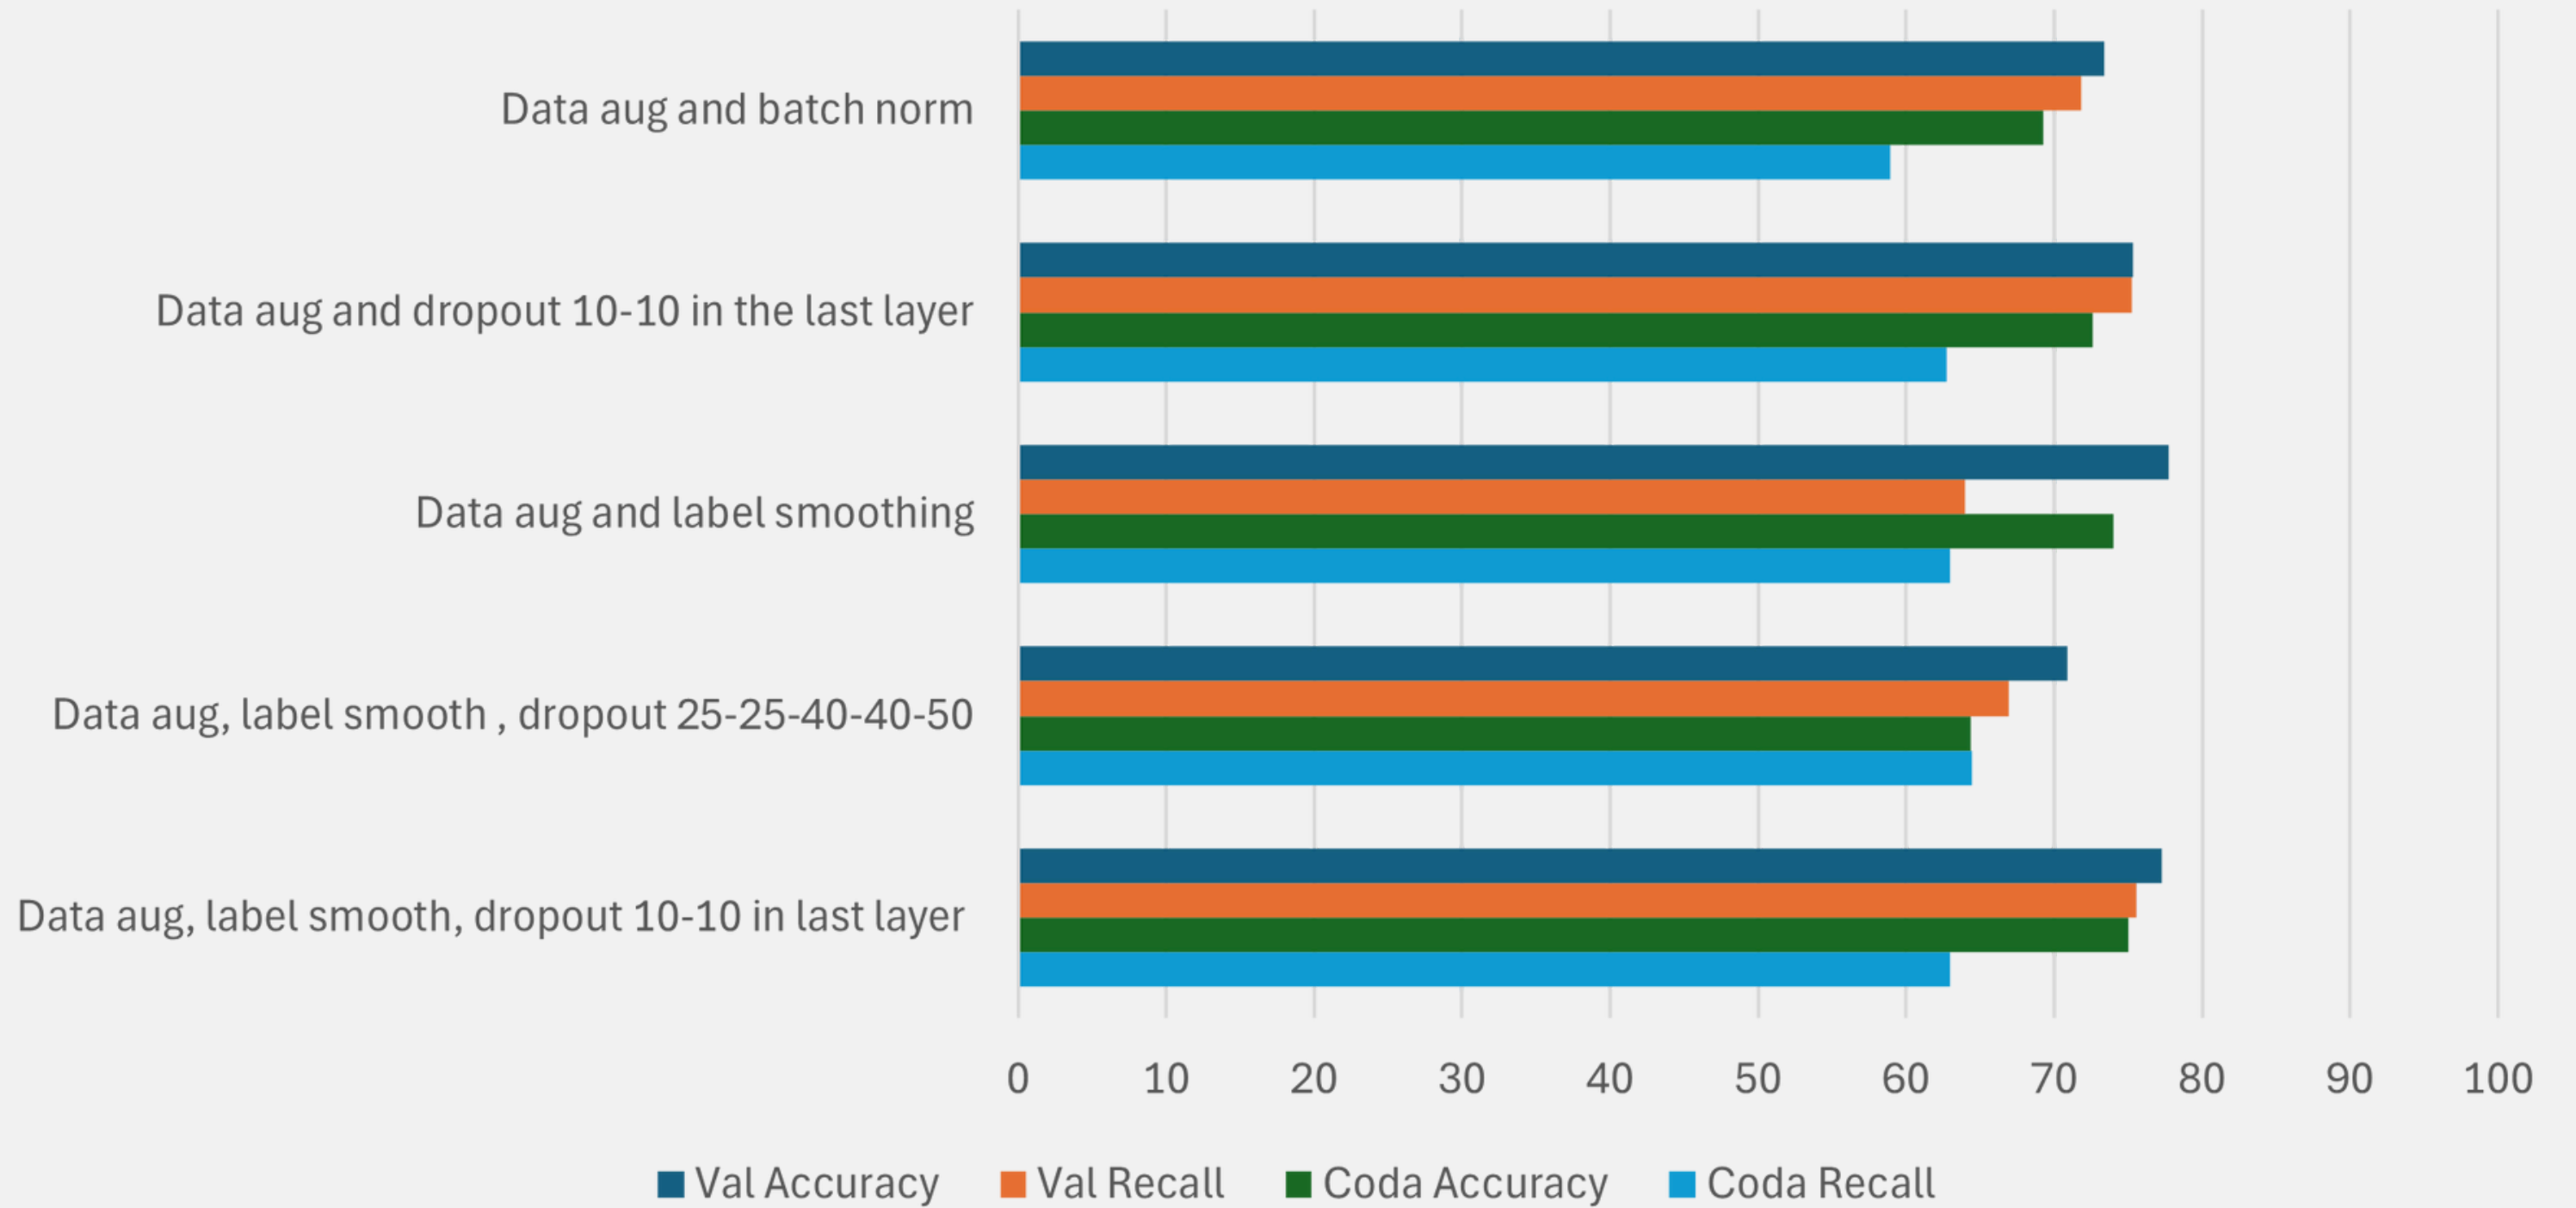
- First used with data augmentation, great results with no overfitting
- Possibly putting more epochs and get even better results

Discussion. Label Smoothing

- Applied to architecture 2 with 20 epochs combined with data augmentation, resulting in a relatively bad performance overall.
- Architecture 5 was taken with 100 epochs and data augmentation as well, performing significantly better, however it tended to overfit.
- A dropout (25–25–40–40–50) was added, which gave worse results than without the dropout, nevertheless it overfitted less.
- Also, a smaller dropout (10) in the dense layers, gave a very similar result with the experiment tried without dropout at all.
- A bigger dropout (30) (+ batch normalization) resulted in a decrease of 15% in validation accuracy.

Discussion. Comparison

Arch 5


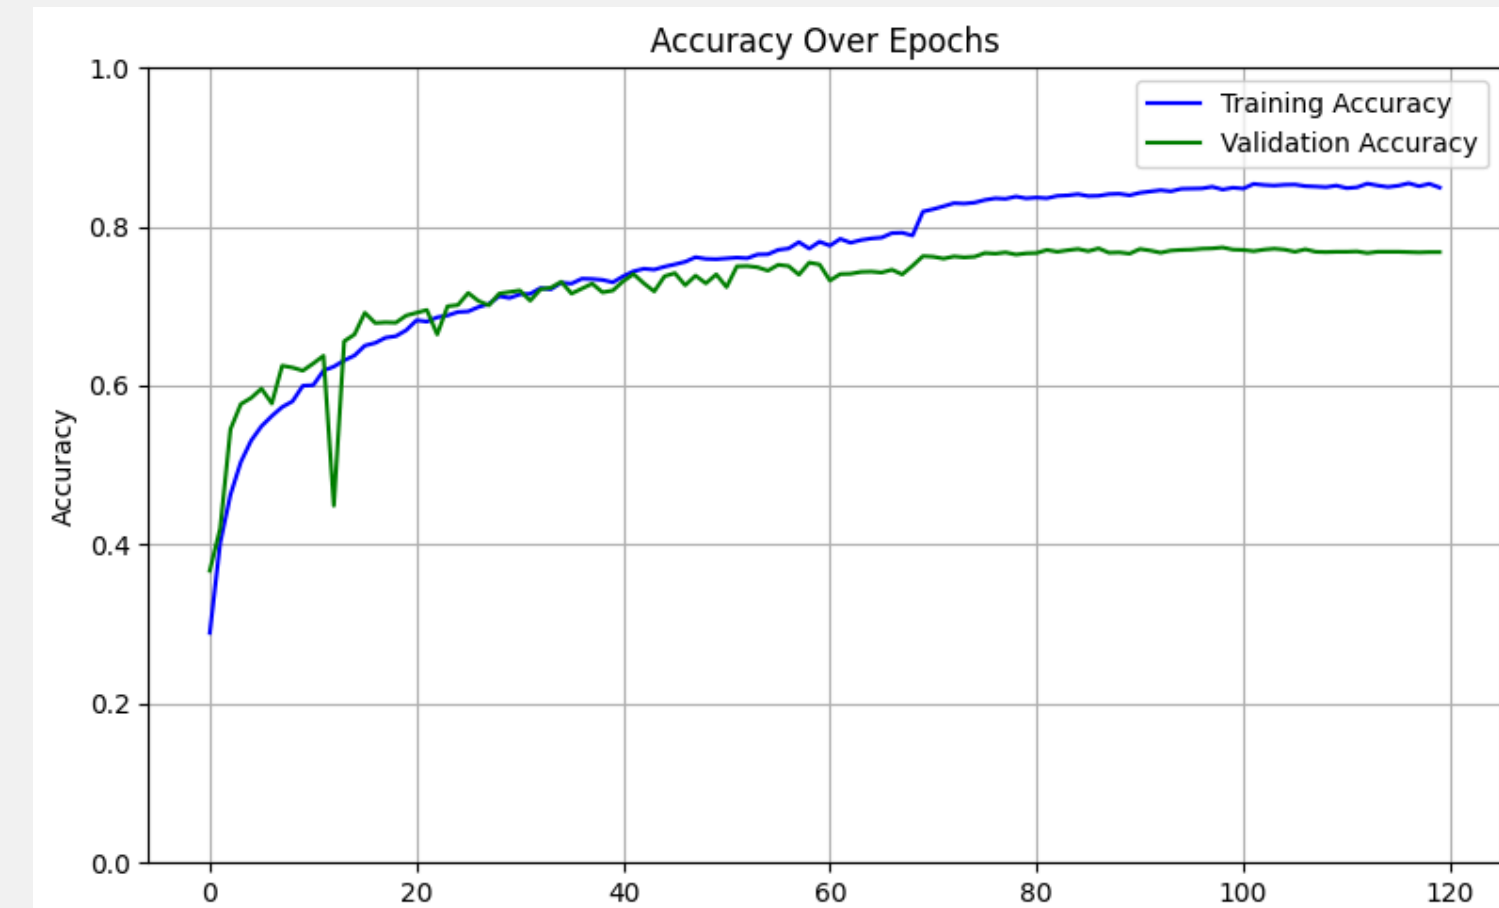


Conclusions

The **best performing neural networks** were:

- network 5 with data augmentation and label smoothing
- network 5 with data augmentation, label smoothing and dropout in the dense layers

Network	Mean Accuracy	Mean Precision	Mean Recall
without dropout	74.72	63.99	63.48
with dropout	75.07	73.87	63.69



Network 5 without dropout

The background is a light gray color, decorated with various hand-drawn blue doodles. These include several loops and swirls at the top, a series of vertical wavy lines on the left, a zigzag line at the bottom center, and several checkmarks or 'v' shapes on the bottom right. There are also some abstract scribbles and lines scattered throughout.

Thank you